

Forecasting Bitcoin Price Movements from Crypto-News Sentiment Analysis

Mahla Entezari

Shahid Beheshti University – Department of Computer Science

June 2024

Abstract

We present an end-to-end machine-learning pipeline that ingests crypto-news articles, quantifies their sentiment, and predicts the *direction* of Bitcoin’s hourly price move immediately after publication. Two public datasets—*Crypto News* (44k articles, 2013–2024) and Binance minute-level OHLCV—are time-aligned, enriched with lexical and market-context features, and fed to four model families (logistic regression, random forest, XGBoost, and a fine-tuned BERT). The best model (Bayesian-optimised XGBoost) achieves 68.4% macro F1 and 71.2% ROC-AUC on an out-of-time test set covering the high-volatility Q1-2024 period. Error analysis reveals sarcasm, information-propagation lag, and macro-economic shocks as main failure drivers. Source-credibility and sentiment polarity contribute most to predictive power.

Contents

1	Introduction	3
2	Problem Formulation	3
2.1	Task Definition	3
2.2	Why Machine Learning?	3
3	Datasets	3
3.1	Crypto News Corpus	3
3.2	Bitcoin OHLCV	4
3.3	Temporal Alignment	4
4	Exploratory Data Analysis	4
4.1	Source Distribution	4
4.2	Sentiment by Subject	5
4.3	Temporal Seasonality	5
5	Feature Engineering	5
6	Modelling Approach	6
7	Evaluation Protocol	6
8	Results	7
9	Error Analysis	7
10	Threats to Validity	7
11	Future Work	7
12	Conclusion	8

1 Introduction

Bitcoin’s price is notoriously sensitive to news narratives: rumours of institutional adoption, regulatory crack-downs, or exchange hacks can move markets within minutes. This study asks:

Can a model, given only the text of a freshly published crypto-news item, anticipate whether Bitcoin will rise, fall, or stay flat in the subsequent hour ?

Answering requires (i) reliable data acquisition, (ii) robust sentiment extraction, (iii) careful feature engineering, (iv) time-series aware evaluation, and (v) critical analysis of model errors and validity threats. The remainder of the article follows this pipeline.

2 Problem Formulation

2.1 Task Definition

We cast the problem as a three-class classification:

$$y_t = \begin{cases} \uparrow & \text{if } r_{t+1} > +0.25\% \\ \rightarrow & \text{if } |r_{t+1}| \leq 0.25\% \\ \downarrow & \text{if } r_{t+1} < -0.25\% \end{cases}$$

where r_{t+1} is the log-return of Bitcoin’s close price from publication minute t to $t + 60$ min.

2.2 Why Machine Learning?

Rule-based heuristics or pure technical indicators ignore nuanced language patterns (e.g. sarcasm). Transformer embeddings and gradient-boosted trees can capture such subtleties while blending structured market data.

3 Datasets

3.1 Crypto News Corpus

Source. Kaggle dataset *Crypto News* (<https://www.kaggle.com/datasets/oliviertvha/crypto-news>). It aggregates 44 938 headlines and summaries from *CoinTelegraph*, *CryptoNews*, and *CryptoPotato* between 2013-02-11 and 2024-04-18. Fields include timestamp (UTC), source, headline, body, and subject tags (*bitcoin*, *altcoin*, etc.).

3.2 Bitcoin OHLCV

Minute-level OHLCV data are fetched from Binance’s REST API (`/api/v3/klines`) and resampled to hourly bars.

3.3 Temporal Alignment

Each article is joined to the *next* hourly price bucket, with a 30-minute buffer to allow information diffusion (Figure ??). Articles outside trading hours do not exist in crypto markets, simplifying alignment.

4 Exploratory Data Analysis

4.1 Source Distribution




Figure 1: **Article count per outlet.** CoinTelegraph contributes roughly 13 k items—over 40 % of the entire corpus—whereas CryptoNews and especially CryptoPotato are much smaller. The grey dashed line shows the global mean count per outlet.

. The dominance of CoinTelegraph raises the risk of *source bias*: if that outlet consistently uses a bullish or bearish tone, the model may confuse “house style” with genuine market signals. We mitigate this by (i) learning a *Source-Credibility Index* (SCI) and (ii) enabling class-weighting in XGBoost so that minority sources

4.2 Sentiment by Subject




Figure 2: **Histogram + KDE of VADER polarity for six topics.** The X-axis is polarity (1 = very negative, +1 = very positive), the Y-axis shows frequency. Bitcoin peaks at about +0.15; Ethereum is centred near 0; NFT exhibits a long negative tail. Translucent curves represent kernel-density estimates.

. Three insights emerge: (1) Bitcoin headlines are mildly skewed positive, matching the overall bullish tone of the period studied. (2) The heavy left tail for NFT reflects 2022’s “crypto winter” and multiple scam reports, suggesting the model should assign extra weight to strongly negative NFT news. (3) Ethereum’s near-symmetric distribution stems from a mix of technical progress (The Merge, L2 roll-outs) and regulatory uncertainty—helpful for reducing variance in the classifier’s predictions.

4.3 Temporal Seasonality

STL decomposition of hourly news counts reveals a weekend trough, whereas price volatility slightly increases. Therefore, binary **Weekend** and cyclic hour-of-day encodings are added.

5 Feature Engineering

F1. Lexical TF-IDF. 1- and 2-gram vectors of lower-cased, lemmatised headlines (10 k dims, ℓ_2 -normalised).

F2. Sentiment Scores. VADER polarity, RoBERTa sentiment [1], TextBlob subjectivity.

- F3. BERT Embeddings.** 768-dim sentence vector from `all-mpnet-base-v2`.
- F4. Source Credibility Index (SCI).** Target-encoding of source on the training split.
- F5. Temporal Features.** Sine/cosine of hour, minutes-since-previous article, rolling news count (1 h window).
- F6. Market Context.** Lagged returns $\Delta p_{t-1}, \Delta p_{t-2}$, Bollinger-band width.

6 Modelling Approach

Four model families are compared:

LogReg Multinomial logistic regression, ℓ_2 penalty, features F1–F2.

RF Random Forest, 300 trees, depth 20.

XGB XGBoost, hyper-parameters tuned by Bayesian search (`max_depth=9`, `eta=0.05`, 600 rounds).

BERT CryptoBERT¹ fine-tuned 3 epochs; logits fused with F4–F6 via a small MLP.

7 Evaluation Protocol

Temporal Split. Train = 2018-01-01 – 2023-06-30 (70 %), Dev = 2023-07-01 – 2023-12-31 (15 %), Test = 2024-01-01 – 2024-04-18 (15 %).

Metrics. Accuracy, macro-F1, and ROC-AUC. Macro-F1 is primary due to class imbalance.

Baselines. *Persistence* (direction = previous hour) and *Majority class* (\rightarrow).

¹<https://huggingface.co/finiteautomata/bert-base-uncased-crypto>

8 Results

Table 1: Performance on the held-out test set (Q1-2024).

Model	Accuracy	Macro-F1	ROC-AUC
Majority	0.424	0.296	0.500
Persistence	0.492	0.333	0.500
LogReg	0.628	0.571	0.643
Random Forest	0.657	0.602	0.679
XGBoost	0.694	0.684	0.712
BERT+MLP	0.673	0.658	0.701

Observations. Sentiment features lift macro-F1 by 4.7 pp over pure TF-IDF. SCI alone adds 1.5 pp ROC-AUC. Transformer embeddings boost recall for \downarrow but triple latency.

9 Error Analysis

Manual inspection of 100 mis-classifications reveals:

- **Sarcasm & Hyperbole.** E.g. “Bitcoin crashes to \$69K” (an ATH) predicted \downarrow .
- **Latency mismatch.** Market sometimes reacts before the 30-min buffer.
- **Macro shocks.** FOMC rate decisions override crypto-specific sentiment.

10 Threats to Validity

Look-ahead bias.

Feature engineering must avoid using future information.

Label noise.

Price moves driven by on-chain factors are unobserved.

Concept drift.

Narratives evolve; scheduled re-training is required.

11 Future Work

1. Integrate on-chain metrics (active addresses, MVRV).
2. Use streaming LLM embeddings (e.g. `text-embedding-3-large`).
3. Extend to ETH, SOL, and alt-basket indices.
4. Back-test RL execution with transaction costs.

12 Conclusion

Combining textual sentiment with minimal price context outperforms naive baselines in forecasting hourly Bitcoin direction. While a macro-F1 of 0.68 is promising, sarcasm handling, latency reduction, and multi-asset generalisation remain open avenues.

References

- [1] Liu, Y. *et al.* (2019). *RoBERTa: A Robustly Optimized BERT Pretraining Approach*. arXiv:1907.11692.
- [2] Olivier Vhalde. (2024). *Crypto News*. Kaggle Dataset. <https://www.kaggle.com/datasets/oliviervha/crypto-news>
- [3] Chen, T. & Guestrin, C. (2016). *XGBoost: A Scalable Tree Boosting System*. Proc. KDD 2016.